



Chartered Financial Data Scientist (CFDS)

Study on the Tail Distribution of Stock Returns: From Empirical Analysis to Factor Explanations

CFDS-11 – Project Presentation – Pierre Nibéron

ALL RIGHTS RESERVED. No part of this work may be reproduced in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage and retrieval system without expressed written permission from the authors.

- Motivation
- Problem Description
- Data
- Whole S&P500 description
- Tails description
- Time clustering
- Sector clustering
- Modeling
 - CAPM
 - Fama-French 3 factors
 - Carhart's
 - Sentiment
- Conclusion

- On Fat-Tails
 - Mention by Dr. Sikandar Siddiqui during CFDS lessons
 - Personal IT Background going into finance
 - Describing and finding explanatory models
 - Reading literature on the subject
 - Using the tools learned during CFDS course
 - Initiation to existing financial models (CAPM, Fama-French, ...)

"large changes tend to be followed by large changes, of either sign, and small changes tend to be followed by small changes." – Benoit B. Mandelbrot, 1963

- First Part : Description and statistical testing
 - Do we observe fat tails which would not exist under normal assumptions ?
Arithmetic vs. Log, Skewness, Kurtosis, JB and visualization
 - What do these tails look like? Are they symmetrical?
99.9th percentile, clustering, t-test, f-test and visualization
- Second Part : Model and explanatory power
 - Do factor models explain extreme returns as well as they explain the full distribution?
CAPM, Fama-French Three factors model, Carhart's Momentum, Sentiment

- S&P500 stock list from wikipedia
- Daily stock prices from 01.01.2010 to 01.05.2026
- yFinance api (data from yahoo finance website)
- Risk free rate, T-Bill 3 month from FRED's website
- Fama-French dataset from panda_datareader library
- News headlines from kaggle (see link in sentiment analysis part)

```
date,Ticker,Open,High,Low,Close,Adj Close,Volume,previous_adj_close,arithmetic return,arithmetic return %,log_return,index,Dividends,Stock Splits,z_score global,extreme
2020-03-09,APA,13.420000076293944,13.699999809265137,9.31999969482422,9.550000190734863,8.224353790283203,28073200.0,17.826614379882812,-0.5386474618778808,-53.86474618778808,-0.7735928034616633,,,,-39.92385295519495,True
2024-04-11,GL,98.4800033569336,98.58000183105467,38.95000076293945,49.16999816894531,48.29851150512695,36577500.0,103.07022857666016,-0.5314019171966409,-53.14019171966409,-0.7580098442455365,,,,-39.1201694230346,True
2019-01-14,PCG,9.210000038146973,9.729999542236328,7.78000020980835,8.380000114440918,8.26025390625,127198800.0,17.338647842407227,-0.5235929594205784,-52.35929594205784,-0.7414826628960451,,,,-38.26778809421307,True
2020-03-09,OXY,15.579999923706056,19.190000534057617,12.039999961853027,12.510000228881836,11.708246231079102,104930300.0,24.39919090270996,-0.5201379308943112,-52.01379308943112,-0.7342565724046701,35302.0,0.79,0.0,-37.89510599646308,True
2025-10-29,FISV,71.36000061035156,76.6500015258789,66.58000183105469,70.5999984741211,70.5999984741211,103454200.0,126.16999816894533,-0.44043750892676065,-44.043750892676066,-0.5806000665408693,,,,-29.97034654228724,True
2014-05-16,TKO,10.550000190734863,11.93000030517578,10.550000190734863,11.270000457763672,9.194170951843262,40965100.0,16.25908088684082,-0.43452086770264464,-43.452086770264465,-0.5700818854377694,,,,-29.427876474593468,True
2019-01-24,PCG,7.900000095367432,14.5,7.300000190734863,13.949999809265137,13.750662803649902,82538300.0,7.875828266143799,0.7459322802606726,74.59322802606727,0.5572886710184197,,,,28.715710908258853,True
2025-08-26,SATS,54.11000061035156,55.18999862670898,50.619998931884766,50.869998931884766,50.869998931884766,46579100.0,29.8799991607666,0.702476585028781,70.2476585028781,0.5320840056598937,,,,27.415792567890204,True
2018-10-04,SMCI,1.5160000324249268,1.5160000324249268,0.8500000238418579,1.2599999904632568,1.2599999904632568,125601000.0,2.140000104904175,-0.41121498659006966,-41.12149865900697,-0.5296941646598585,,,,-27.34489944561519,True
2024-07-26,DXCM,66.0,67.44000244140625,62.34000015258789,64.0,64.0,53914100.0,107.8499984741211,-0.4065832090358631,-40.65832090358631,-0.5218582753273351,,,,-26.940767269624978,True
2025-07-02,CNC,37.34000015258789,37.77999877929688,33.7599983215332,33.779998779296875,33.779998779296875,93566700.0,56.650001525878906,-0.40370701024843814,-40.37070102484381,-0.517023139166418,,,,-26.69139747699531,True
2015-04-13,BLDR,11.0,11.739999771118164,10.5,11.56999969482422,11.56999969482422,17377500.0,6.900000095367432,0.6768115267986974,67.68115267986974,0.5168940894045324,,,,26.63238003500654,True
```

- Arithmetic>Returns stats table:

```
count      1710528.000000
mean        0.000691
std         0.019362
variance    0.000375
min         -0.538647
1%          -0.052116
5%          -0.027356
25%         -0.007924
median      0.000739
75%         0.009405
95%         0.028037
99%         0.053622
max         0.745932
skewness    0.180127
kurtosis    23.823084
```

- Jarque-Berra

JB statistic: 40458666.05

p-value: 0.00e+00

- Log-Returns stats table:
(check skewness and max,
effect of the log)


```
count      1710528.000000
mean        0.000503
std         0.019403
variance    0.000376
min         -0.773593
1%          -0.053523
5%          -0.027737
25%         -0.007955
median      0.000738
75%         0.009361
95%         0.027652
99%         0.052233
max         0.557289
skewness    -0.580858
kurtosis    26.745270
```

- Jarque-Berra


JB statistic: 51077412.92

p-value: 0.00e+00


Distribution of daily log returns (linear scale)




Daily log returns (log density scale)




QQ-plot of standardized log returns vs Normal






Log-returns vs Normal distribution





Whole set visualization





- Do the two top 0.1% have the same statistical resemblance?
Check both return definition, log_return (left), arithmetic_return (right)
Extreme Value Theory and Generalized Pareto Distribution territory.

```
# is there unique stock contributing?
extreme_returns["Ticker"].nunique()

[5]
... 342
```

$$r_{arithmetic} = e^{r_{log}} - 1$$

$$e^{0.557} - 1 = 0.746$$

$$e^{-0.773} - 1 = -0.539$$

	positive_tail	negative_tail_abs
count	692.000000	1019.000000
mean	0.183787	0.197229
std	0.049444	0.068200
min	0.143244	0.143210
25%	0.153815	0.156726
50%	0.167282	0.176181
75%	0.197947	0.210945
max	0.557289	0.773593

	test	t_stat	p_value
0	Welch t-test	-4.723884	0.000003
1	Levene test	13.545443	0.000240

	positive_tail	negative_tail_abs
count	692.000000	1019.000000
mean	0.203309	0.177224
std	0.064695	0.050491
min	0.154012	0.133428
25%	0.166275	0.145062
50%	0.182088	0.161534
75%	0.218897	0.190182
max	0.745932	0.538647

	test	t_stat	p_value
0	Welch t-test	8.920870	1.628379e-18
1	Levene test	5.170664	2.309497e-02

- Top 0.1% daily returns clustered by years

year	
2010	39
2011	73
2012	65
2013	51
2014	24
2015	53
2016	56
2017	37
2018	54
2019	61
2020	763
2021	20
2022	69
2023	51
2024	97
2025	176
2026	44



- Top 0.1% daily returns clustered by days, ordered, head(10)

```
Date
2020-03-16    185
2020-03-24     92
2020-03-18     80
2020-03-12     79
2020-03-09     46
2020-11-09     36
2020-03-13     35
2025-04-09     30
2020-03-19     25
2020-03-17     23
```


- What years were good / bad for the top 0.1%?

direction	negative_event	positive_event	net	ratio_positive	mean_log_return	median_log_return
year						
2010	17	22	5	0.564103	0.014507	0.146681
2011	44	28	-16	0.388889	-0.050225	-0.149092
2012	39	26	-13	0.400000	-0.038452	-0.153557
2013	21	29	8	0.580000	0.040606	0.153643
2014	10	14	4	0.583333	0.014659	0.151223
2015	29	24	-5	0.452830	-0.018127	-0.147640
2016	34	22	-12	0.392857	-0.044785	-0.153643
2017	20	17	-3	0.459459	-0.021121	-0.144722
2018	32	21	-11	0.396226	-0.048438	-0.155941
2019	35	26	-9	0.426230	-0.049654	-0.150451
2020	470	280	-190	0.373333	-0.053402	-0.153259
2021	10	10	0	0.500000	-0.015867	-0.000785
2022	49	20	-29	0.289855	-0.081468	-0.151663
2023	31	19	-12	0.380000	-0.064066	-0.160101
2024	52	44	-8	0.458333	-0.031347	-0.146023
2025	96	76	-20	0.441860	-0.031597	-0.150361
2026	30	14	-16	0.318182	-0.071460	-0.155533

- What does it look like? (Arithmetic returns to not crush the extremes positives)



- What does it look like? (Arithmetic returns to not crush the extremes positives)



- Do the top 0.1% mean across years are statistically different within each positive extreme and negative extreme? log_return (left), arithmetic_return (right)

Less impact of return definition because we are comparing positive / positive and negative / negative.

Extreme Value Theory and Generalized Pareto Distribution territory.

tail in log return				tail in arithmetic return			
		Welch-ANOVA	p_value			Welch-ANOVA	p_value
0	positive_extremes_across_years	1.387477	0.158121	0	positive_extremes_across_years	1.461559	0.124434
1	negative_extremes_across_years	0.924244	0.543047	1	negative_extremes_across_years	0.957364	0.506226

- Top 0.1% daily returns clustered by sectors, ordered
Raw count and normalized (events/nbr_stocks) count.



sector	
Information Technology	410
Consumer Discretionary	274
Industrials	198
Health Care	175
Financials	172
Energy	118
Communication Services	88
Real Estate	77
Utilities	75
Materials	74
Consumer Staples	72

sector	
Consumer Discretionary	5.708333
Information Technology	5.616438
Energy	5.363636
Communication Services	3.826087
Health Care	3.017241
Materials	2.846154
Industrials	2.506329
Real Estate	2.483871
Utilities	2.419355
Financials	2.263158
Consumer Staples	2.000000


- What sectors performed best in the top 0.1%? Ordered table.

	mean_log_return	median_log_return	positive_ratio
sector			
Communication Services	0.0054	0.1443	0.5581
Information Technology	-0.0251	-0.1477	0.4496
Energy	-0.0333	-0.1482	0.4609
Industrials	-0.0424	-0.1506	0.3776
Materials	-0.0434	-0.1513	0.3973
Consumer Discretionary	-0.0447	-0.1521	0.4037
Health Care	-0.0533	-0.1547	0.3829
Consumer Staples	-0.0664	-0.1562	0.3239
Financials	-0.0678	-0.1535	0.3412
Utilities	-0.0719	-0.1581	0.3333
Real Estate	-0.0782	-0.1656	0.3158

- What does it look like? (Arithmetic returns to not crush the extremes positives)



- What does it look like? (Arithmetic returns to not crush the extremes positives)



- Do the top 0.1% mean across sectors are statistically different within each positive extreme and negative extreme? log_return (left), arithmetic_return (right)

*Less impact of return definition because we are comparing positive / positive and negative / negative.
Extreme Value Theory and Generalized Pareto Distribution territory.*

	tail in log return	Welch-ANOVA	p_value		tail in arithmetic return	Welch-ANOVA	p_value
0	positive_extremes_across_sectors	2.607491	0.005759	0	positive_extremes_across_sectors	2.683449	0.004538
1	negative_extremes_across_sectors	3.426325	0.000301	1	negative_extremes_across_sectors	3.445695	0.000281

- Do factor models explain extreme returns as well as they explain the full distribution?

$$\text{CAPM : } R_i - R_f = \alpha + \beta(R_m - R_f)$$

R_i — the return of your stock (or portfolio)

R_m — the return of the market (S&P500 index)

R_f — the risk-free rate (DTB3, the 3-month T-bill)

Build a simple OLS for CAPM on AAPL to test.

```

=====
                        OLS Regression Results
=====
Dep. Variable:          stock_excess      R-squared:                0.470
Model:                  OLS              Adj. R-squared:          0.470
Method:                 Least Squares     F-statistic:             3485.
Date:                  Sun, 03 May 2026   Prob (F-statistic):       0.00
Time:                  15:10:02          Log-Likelihood:          11494.
No. Observations:      3936             AIC:                    -2.298e+04
Df Residuals:          3934             BIC:                    -2.297e+04
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0005	0.000	2.629	0.009	0.000	0.001
market_excess	1.1211	0.019	59.036	0.000	1.084	1.158

```

=====
Omnibus:                551.269      Durbin-Watson:           1.908
Prob(Omnibus):          0.000        Jarque-Bera (JB):        7574.137
Skew:                   0.075        Prob(JB):                0.00
Kurtosis:               9.794        Cond. No.                 91.3
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```



Beta = 1.12 and highly significant ($t = 59$, $p \approx 0$). AAPL moves slightly more than the market, a 1% market move corresponds to a 1.12% AAPL move on average.

$R^2 = 0.47$ means 47% of AAPL's daily return variance is explained by the market alone. For a single factor model on a single stock, that is respectable. The remaining 53% is idiosyncratic, company-specific news, earnings, product launches etc.

Alpha = 0.0005 and significant ($p = 0.009$). This is small but tells you AAPL generates a tiny daily excess return above what CAPM predicts, annualized that is roughly $0.0005 \times 252 \approx 12.6\%$ annual alpha.

Jarque-Bera = 7574, $p \approx 0$, the residuals are overwhelmingly non-normal.

Kurtosis = 9.79, the residuals have fat tails.



Do the CAPM regression on each individual stock for the whole set and summarize.

	Ticker	alpha	beta	r2	p_alpha	p_beta	n
0	A	0.000134	1.157895	0.480006	0.523352	0.0	3936
1	AAPL	0.000547	1.121092	0.469760	0.008605	0.0	3936
2	ABT	0.000246	0.738777	0.348994	0.163096	0.0	3927
3	ACGL	0.000352	0.793867	0.334955	0.069585	0.0	3989
4	ACN	0.000133	1.021685	0.515069	0.442290	0.0	3949

	metric	mean	median	std	min	max
0	alpha	0.000201	0.000191	0.000237	-0.000457	0.001104
1	beta	0.980084	0.981246	0.283460	0.356161	1.694702
2	r2	0.337979	0.336831	0.118038	0.026121	0.642840

Beta Median beta = 0.98, almost exactly 1.0. This is completely expected with S&P500 constituents, which by construction are the market.

Median $R^2 = 0.337$ means the market factor explains about 34% of individual stock return variance on average.


Alpha

Only 9.2% of stocks show significant alpha at the 5% level.


The mean alpha = 0.0002 per day, which annualizes to roughly $0.0002 \times 252 \approx 5\%$.

CAPM Full Sample — Cross-Sectional Distribution (423 stocks)


Distribution of R^2 — Full Sample



Distribution of Beta — Full Sample



Distribution of Annualized Alpha — Full Sample



Now test the CAPM on the top 0.1%.

We first try with one ticker ("SMCI") before looping and aggregating.

	full_tail	pos_tail	neg_tail
Ticker			
SMCI	47	26.0	21.0
PCG	24	8.0	16.0
EXPE	23	11.0	12.0
AXON	22	13.0	9.0
APA	22	10.0	12.0
CCL	22	9.0	13.0
RCL	22	12.0	10.0
AMD	22	10.0	12.0
COHR	20	11.0	9.0
FSLR	20	12.0	8.0

Now test the CAPM on the top 0.1%.

We first try with one ticker ("SMCI") before looping and aggregating.

```
=====
SMCI - FULL SAMPLE
=====
                        OLS Regression Results
=====
Dep. Variable:          stock_excess      R-squared:                0.165
Model:                  OLS              Adj. R-squared:          0.165
Method:                 Least Squares     F-statistic:             788.5
Date:                   Sat, 09 May 2026   Prob (F-statistic):      1.74e-158
Time:                   13:30:20          Log-Likelihood:          7668.0
No. Observations:       3991             AIC:                    -1.533e+04
Df Residuals:           3989             BIC:                    -1.532e+04
Df Model:                1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0008	0.001	1.494	0.135	-0.000	0.002
market_excess	1.4423	0.051	28.081	0.000	1.342	1.543

```
=====
```


Now test the CAPM on the top 0.1%.

We first try with one ticker ("SMCI") before looping and aggregating.


Doing the OLS on the full tail is actually a mistake. (Mixing artifacts.)

```
=====
SMCI - FULL TAIL
=====
                        OLS Regression Results
=====
Dep. Variable:          stock_excess      R-squared:                0.144
Model:                  OLS              Adj. R-squared:          0.125
Method:                 Least Squares     F-statistic:             7.591
Date:                  Sat, 09 May 2026   Prob (F-statistic):      0.00844
Time:                  13:30:20          Log-Likelihood:          8.1068
No. Observations:      47               AIC:                    -12.21
Df Residuals:          45               BIC:                    -8.513
Df Model:              1
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                0.0205      0.030      0.673      0.504     -0.041      0.082
market_excess        3.2707      1.187      2.755      0.008      0.880      5.662
```

Now test the CAPM on the top 0.1%.

We first try with one ticker ("SMCI") before looping and aggregating.

Doing the OLS on the full tail is actually a mistake. (Mixing artifacts.)



Each cloud is flat — the steep line only exists because the two clouds are vertically separated

Now test the CAPM on the top 0.1%.

We first try with one ticker ("SMCI") before looping and aggregating.

```
=====
SMCI - NEGATIVE TAIL
=====
                        OLS Regression Results
=====
Dep. Variable:          stock_excess    R-squared:                0.016
Model:                  OLS             Adj. R-squared:          -0.036
Method:                 Least Squares    F-statistic:              0.3096
Date:                   Sat, 09 May 2026  Prob (F-statistic):      0.584
Time:                   13:30:20         Log-Likelihood:           27.504
No. Observations:       21              AIC:                    -51.01
Df Residuals:           19              BIC:                    -48.92
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.2047	0.016	-13.063	0.000	-0.238	-0.172
market_excess	0.2928	0.526	0.556	0.584	-0.809	1.394

```
=====
```

```
=====
SMCI - POSITIVE TAIL
=====
                        OLS Regression Results
=====
Dep. Variable:          stock_excess    R-squared:                0.000
Model:                  OLS             Adj. R-squared:          -0.041
Method:                 Least Squares    F-statistic:              0.004744
Date:                   Sat, 09 May 2026  Prob (F-statistic):      0.946
Time:                   13:30:20         Log-Likelihood:           35.435
No. Observations:       26              AIC:                    -66.87
Df Residuals:           24              BIC:                    -64.35
Df Model:                1
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.2171	0.015	14.927	0.000	0.187	0.247
market_excess	-0.0460	0.668	-0.069	0.946	-1.426	1.334

```
=====
```

Now test the CAPM on the top 0.1%.

Here, separating on both tails. Minimum of 5 extreme obs. per ticker for the OLS.

	sample	n_stocks	median_r2	mean_r2	median_beta	median_alpha_ann	pct_sig_alpha	pct_sig_beta
0	Full Sample	423	0.337	0.338	0.981	0.048	9.220	100.000
1	Positive Tail	40	0.094	0.171	-0.065	50.902	100.000	7.500
2	Negative Tail	61	0.129	0.212	0.064	-43.149	95.082	11.475

pct_sig_alpha & pct_sig_beta = percentage of stock where alpha or beta had a significant p-value (>0.05).

CAPM shows very little correlation between stock excess and market excess in the tails !

- Do factor models explain extreme returns as well as they explain the full distribution?

$$FF3 : R_i - R_f = \alpha + \beta(R_m - R_f) + \beta_{SMB} \cdot SMB + \beta_{HML} \cdot HML$$

R_i - the return of your stock (or portfolio)

R_m - the return of the market (S&P500 index)

R_f - the risk-free rate (DTB3, the 3-month T-bill)

SMB - Small minus Big (market cap factors)

HML - High minus Low (book to market ratio)

- FF3 on the whole set of S&P500 from 01.01.2010

```

CAPM median adjusted R² : 0.337
FF3  median adjusted R² : 0.373
Gain                               : 0.036

                                0
sample          FF3 – Full Sample
n_stocks                423
median_r2              0.373295
median_adj_r2          0.372834
mean_r2                0.380349
mean_adj_r2            0.379894
median_beta_mkt         0.9413
median_beta_smb         0.062804
median_beta_hml         0.207891
median_alpha_ann        0.03382
pct_sig_alpha           6.382979
pct_sig_smb             81.08747
pct_sig_hml            90.543735

```


- FF3 for the top 0.1%. R2 Increases ! ... but no, it is overfit.

	0	1	2
sample	FF3 – Full Sample	FF3 – Positive Tail	FF3 – Negative Tail
median_r2	0.373295	0.520654	0.581733
median_beta_mkt	0.9413	0.216505	0.23334
median_beta_smb	0.062804	-0.311619	-0.006625
median_beta_hml	0.207891	-0.032717	-0.089737
median_alpha_ann	0.03382	49.177992	-42.534714
pct_sig_alpha	6.382979	85.714286	75.409836
pct_sig_smb	81.08747	7.142857	1.639344
pct_sig_hml	90.543735	4.761905	4.918033

- FF3 for the top 0.1%. With adjusted R2.

sample	FF3 – Full Sample	FF3 – Positive Tail	FF3 – Negative Tail
median_r2	0.373295	0.520654	0.581733
median_adj_r2	0.372834	0.072568	0.156952
median_beta_mkt	0.9413	0.216505	0.23334
median_beta_smb	0.062804	-0.311619	-0.006625
median_beta_hml	0.207891	-0.032717	-0.089737
median_alpha_ann	0.03382	49.177992	-42.534714
pct_sig_alpha	6.382979	85.714286	75.409836
pct_sig_smb	81.08747	7.142857	1.639344
pct_sig_hml	90.543735	4.761905	4.918033

Neither CAPM nor FF3 seems good to explain extreme returns. Adding size and value factors absorbs some alpha but the vast majority of tail behavior remains idiosyncratic and unexplained by systematic factors.

- Do extreme returns have a persistence component to them?
The case for momentum also known as the Carhart 4 factors model.

$$R_i - R_f = \alpha + \beta_m(R_m - R_f) + \beta_{SMB} \cdot SMB + \beta_{HML} \cdot HML + \beta_{UMD} \cdot UMD$$

R_i - the return of your stock (or portfolio)

R_m - the return of the market (S&P500 index)

R_f - the risk-free rate (DTB3, the 3-month T-bill)

SMB - Small minus Big (market cap factors)

HML - High minus Low (book to market ratio)

UMD - Up minus Down (long time winners vs long time losers)

Adjusted R2 doubles in the tails and the percentage of stocks with a significant alpha is halved ! Rallies go against momentum. Crashes follow it. The UMD significance tells us that it explains crashes better than rallies.

	0	1	2
sample	Carhart — Full Sample	Carhart — Positive Tail	Carhart — Negative Tail
n_stocks	423	42	61
median_r2	0.376237	0.692238	0.854992
median_adj_r2	0.375625	0.137812	0.289059
median_beta_mkt	0.939983	0.522633	0.158503
median_beta_smb	0.053771	-0.085258	0.601842
median_beta_hml	0.196626	0.001649	0.34015
median_beta_umd	-0.032806	-0.292279	1.104322
median_alpha_ann	0.035172	47.957742	-41.121436
pct_sig_alpha	6.146572	50.0	62.295082
pct_sig_umd	59.810875	2.380952	11.47541

Does market sentiment explain these extreme returns?
(using a proxy of financial news headlines)

$$R_i - R_f = \alpha + \beta_m(R_m - R_f) + \beta_{SMB} \cdot SMB + \beta_{HML} \cdot HML + \beta_{UMD} \cdot UMD + \beta_{SENT} \cdot SENT$$

R_i - the return of your stock (or portfolio)

R_m - the return of the market (S&P500 index)

R_f - the risk-free rate (DTB3, the 3-month T-bill)

SMB - Small minus Big (market cap factors)

HML - High minus Low (book to market ratio)

UMD - Up minus Down (long time winners vs long time losers)

SENT - minus 1 for low sentiment, plus 1 for high sentiment


Does market sentiment explain these extreme returns?
(using a proxy of financial news headlines)

Kaggle S&P 500 with Financial News Headlines (2008–2024)

▲ Title	📅 Date	# CP
JPMorgan Predicts 2008 Will Be "Nothing But Net"	2008-01-02	1447.16
Dow Tallies Biggest First-session-of-year Point Drop Ever	2008-01-02	1447.16
2008 predictions for the S&P 500	2008-01-02	1447.16
U.S. Stocks Higher After Economic Data, Monsanto Outlook	2008-01-03	1447.16
U.S. Stocks Climb As Hopes Increase For More Fed Moves	2008-01-07	1416.18

Does market sentiment explain these extreme returns?

But limitations of the dataset... (sparsity and cutoff)



Does market sentiment explain these extreme returns?

FinBERT for NLP analysis :

Hugging Face FinBERT

FinBERT is a pre-trained NLP model to analyze sentiment of financial text. It is built by further training the BERT language model in the finance domain, using a large financial corpus and thereby fine-tuning it for financial sentiment classification. Financial PhraseBank by Malo et al. (2014) is used for fine-tuning.

```
Sentiment distribution:  
sentiment_label  
neutral      10816  
negative     4690  
positive     3367
```

Does market sentiment explain these extreme returns?

	0	1	2
sample	Sentiment — Full Sample	Sentiment — Positive Tail	Sentiment — Negative Tail
n_stocks	423	32	47
median_r2	0.413318	0.904962	1.0
median_adj_r2	0.41243	0.259324	0.291162
median_beta_mkt	0.964401	-0.066094	0.087291
median_beta_smb	0.040218	-0.162851	1.164095
median_beta_hml	0.171875	-0.640028	0.197366
median_beta_umd	-0.00604	-0.3372	0.993179
median_beta_sentiment	-0.000361	-0.042487	-0.012072
median_alpha_ann	0.032603	46.67703	-38.152934
pct_sig_alpha	7.565012	31.25	27.659574
pct_sig_sentiment	5.437352	0.0	4.255319

Does market sentiment explain these extreme returns?

	alpha	beta_mkt	beta_smb	beta_hml	beta_umd	beta_sentiment
alpha	1.000000	0.092686	0.290206	-0.394858	0.425370	0.127724
beta_mkt	0.092686	1.000000	0.600818	0.152756	-0.415668	0.174555
beta_smb	0.290206	0.600818	1.000000	0.106909	-0.331364	0.277959
beta_hml	-0.394858	0.152756	0.106909	1.000000	-0.449070	0.057881
beta_umd	0.425370	-0.415668	-0.331364	-0.449070	1.000000	-0.341020
beta_sentiment	0.127724	0.174555	0.277959	0.057881	-0.341020	1.000000

	Model	Whole Set Median Adj R ²	Positive Tail Median Adj R ²	Negative Tail Median Adj R ²
0	CAPM	0.34	0.09	0.13
1	FF3	0.37	0.07	0.15
2	MOM	0.37	0.14	0.29
3	Sentiment	0.41	0.25	0.29

- 1) Overall returns on the S&P500 do exhibit fat tail behavior.
- 2) Definition of returns, log vs. arithmetic, starts to really matter in the tails.
- 3) Positive tail and negative tail are not symmetric. Neither in count nor magnitude.
- 4) Extremes distribution seem to be consistent across years.
- 5) Extremes distribution do not seem to be consistent across sectors. (Probably due to the Communication Services outlier. To be tested.)
- 6) Modeling in the tails performs worse than on the full set of returns.
- 7) Factors have very low statistical significance ($p < 0.05$) in the tails overall.
- 8) Crashes, negative tail, seems to be best explained using CAPM and momentum.
- 9) Rallies, positive tail, seems to have very shaky significance with the models used overall.
- 10) I learned about extreme value theory and generalized pareto distributions while making this presentation. So, there is an exploration of these tools to be done next time.



Questions?